

THE MATHEMATICS OF

THE THIRD TRADE

*An Expected Value Study of Hit Rate, Ratio, Cost, and Conditional
Trade-Stacking in BankNifty Options*

A Working Notebook

A Personal Trading Research Notebook

July 2026 - Third Edition

Contents

Preface	The Question Behind the Numbers
Chapter 1	Real-World Proof First — The FXCM Study
Chapter 2	The Shape of a Single Trade
Chapter 3	Why Traders Chase Hit Rate
Chapter 4	The Ceiling — Why 50% Is Already Near the Top
Chapter 5	What Actually Moves the Needle: Entries, Exits, Ratio
Chapter 6	Three Ways to Stop — The Conditional Ladder
Chapter 7	What Happens When You Never Stop
Chapter 8	A Softer Edge — Testing a 1.5:1 Ratio
Chapter 9	Returning to 2:1, and Counting by the Month
Chapter 10	Lot Size and the Cost Floor
Chapter 11	The Middle Path — A Conditional Third Trade
Chapter 12	The Argument for No Ceiling
Chapter 13	Three Setups a Day — The Chosen Plan
Chapter 14	Guarding Against the Day the Math Breaks
Closing Note	What This Notebook Is, and Isn't
Appendix	Full Reference Tables

Preface

The Question Behind the Numbers

Every trading plan is, underneath its charts and setups, a small piece of arithmetic repeated many times. A trader risks a fixed amount to make a multiple of it, does this with some probability of success, and repeats the process until the pattern either compounds into something durable or erodes under its own variance. The hard part is rarely the arithmetic. The hard part is deciding, in advance, how many times a day to repeat it, what ratio to demand, and under what conditions to stop.

This notebook grew out of exactly that question, asked about BankNifty options trading. Rather than starting from an arbitrary assumption, it opens with real evidence about how retail traders actually perform, because that evidence is what justifies every hit-rate assumption made afterward. From there, it works through a fixed profit-to-loss ratio, a fixed hit rate, a flat per-trade cost, and a simple stopping rule, before widening into what a trader should actually be trying to improve — the hit rate, or the ratio — and what lot sizing does to the real, after-cost edge.

A Note on Timing: Read This Before You Trade, Not After

There is a reason this notebook is worth reading before a first trade is ever placed, rather than after the hundredth. Almost everything in the chapters ahead is unlearning: the instinct to chase a higher win rate, the belief that being right often enough is the goal, the habit of closing for the day the moment a trade goes well. **These habits form quickly, usually within the first few weeks of live trading, and they are far harder to undo once real money and real ego are both attached to them.** A trader who has never placed a trade has, in one sense, an advantage over a trader with a year of screen time behind them: nothing yet to unlearn. This notebook is written to make the most of that advantage — to install the correct instinct, ratio over hit rate, arithmetic over feeling, before the wrong one ever gets the chance to take hold. A trader already active in the markets will still find every argument here fully intact; it simply arrives as a correction rather than a foundation.

What follows is not investment advice, and it does not know what tomorrow's market will do. It is a record of a calculation — clean, assumption-bound, and useful precisely because its assumptions are stated rather than hidden, and grounded, from the first page, in what real trading data actually shows.

“The edge lives in the trade. The risk lives in the sequence.”

Chapter One

Real-World Proof First — The FXCM Study

Before any spreadsheet arithmetic, this notebook starts with a question worth asking honestly: when real traders are right more often than not, do they actually make money? A broker-scale study of live retail trading gives a direct, well-documented answer, and that answer shapes every assumption used from here forward — most importantly, why this notebook holds hit rate at a realistic 50% throughout, rather than assuming a higher number simply because it would be more comfortable to work with.

FXCM, a retail forex broker, published research examining millions of live client trades across major currency pairs, later followed by a second analysis covering an even larger set of trades. The shape of the finding was consistent across both: retail traders were, on average, right more often than not, closing a majority of their trades at a profit. And yet, as a group, they lost money.

The Core Finding

Looking specifically at EUR/USD, the research found that trades closed profitably a little over three times out of five. But the average size of a losing trade ran noticeably larger than the average size of a winning one — on the order of 70% larger. Traders were, in effect, winning the majority of their battles while losing the war, because each individual loss quietly outweighed each individual win.

Metric	Value
Trades closing at a profit (EUR/USD)	≈61%
Average winning trade	≈48 pips
Average losing trade	≈83 pips
Loss size relative to win size	≈70% larger

Table 1.1 – Summary of the FXCM EUR/USD finding.

EXAMPLE — Doing the EV math on the FXCM numbers ourselves

*Using the same expected-value approach this entire notebook relies on: $EV = (0.61 \times 48) - (0.39 \times 83) = 29.28 - 32.37 = -3.09$ **pips per trade, on average.** A trader winning three trades out of five was still, on average, bleeding value with every trade taken — not because the calling of direction was wrong more often than right, but because the size of being wrong was structurally larger than the size of being right.*

What Separated the Traders Who Actually Profited

The research went further and split traders into two groups: those who maintained a ratio of at least one unit of profit for every unit risked, and those who didn't. The difference in outcomes was large. **Traders who held a ratio of 1:1 or better turned a profit roughly three times as often as those who didn't** — a gap driven purely by ratio discipline, not by any difference in how often either group was right about market direction.

The research's own practical advice mirrors the logic built up across this notebook: cut a losing trade early rather than let it grow, and give a winning trade room to extend, so that the size of a typical win comes to exceed the size of a typical loss — a ratio fix, not a prediction fix.

Why This Notebook Holds Hit Rate at 50% Throughout

This is the single most important thing this study establishes for everything that follows. The FXCM traders were not unlucky guessers — they were often right more than half the time, in some cases well above 60%, and it still was not enough. **A higher hit rate did not protect them, because the ratio they were running was working against them at the same time.**

That is precisely why every calculation in this notebook holds the hit rate at a plain, unglamorous 50% rather than assuming something more flattering like 60% or 70%. Assuming a higher hit rate would have been easy — it would have made every later table look more generous. But it would also have quietly repeated the exact mistake this study documents: leaning on hit rate as the source of comfort, when the real evidence shows hit rate alone guarantees nothing. **Holding the assumption at 50% forces the rest of this notebook to prove its case using the ratio and the structure of the trades themselves** — the two levers the FXCM data shows actually matter — rather than quietly borrowing false comfort from an optimistic win rate.

Every chapter from here forward inherits this discipline. Wherever a hit rate is needed for a calculation, this notebook uses 50%, not because it is the best a trader could hope for, but because it is a realistic, defensible number that does not let the rest of the analysis cheat.

“Being right most of the time is not the same as making money most of the time — that is why this notebook refuses to assume more than a coin-toss-level hit rate, and asks the ratio to do the real work instead.”

Chapter Two

The Shape of a Single Trade

With the reasoning behind a realistic hit rate established, the notebook can now build its own model. Before any daily structure can be judged, the single trade must be understood on its own terms. The founding assumptions were these: a profit-to-loss ratio of 1:2, meaning a winning trade returns ₹1,000 and a losing trade costs ₹500; a hit rate held flat at 50%, for the reasons just laid out in Chapter One; and a flat overhead of ₹100 per trade for platform fees, brokerage, and slippage — charged whether the trade wins or loses.

Outcome	Gross Result	Cost	Net Result
Win	₹1,000	-₹100	+₹900
Loss	-₹500	-₹100	-₹600

Table 2.1 – Net result of a single trade after cost.

$$EV(\text{one trade}) = 0.5 \times (+₹900) + 0.5 \times (-₹600) = +₹150$$

EXAMPLE — One trade, ten times over

Suppose this exact trade is taken ten times, independently, at a true 50% hit rate. On average, five will win and five will lose: $5 \times ₹900 = ₹4,500$ from winners, and $5 \times ₹600 = ₹3,000$ lost on losers. Net: +₹1,500 across ten trades, or +₹150 per trade — exactly the EV figure above.

This single number is the quiet engine of everything that follows. Every daily structure examined in this notebook — one trade, two trades, three trades, conditional or unconditional — is ultimately built from repetitions of this same ₹150 edge.

Chapter Three

Why Traders Chase Hit Rate

Chapter One showed, with real data, that a high hit rate can coexist with a losing system. It is worth pausing to ask why traders chase hit rate anyway, because the fixation is understandable even when it is misplaced.

Hit rate is the most emotionally visible number in trading. Every single trade closes with an immediate, binary verdict — win or lose — and that verdict arrives instantly and personally. A trader who loses six of the last ten trades feels that string of losses acutely, one at a time, regardless of what the sizing on those trades was. The profit ratio, by contrast, only pays out its true meaning in aggregate, over many trades — it is an abstraction that rewards patience, while hit rate is a scoreboard that updates after every single point.

This creates a natural pull toward systems that are engineered to win often — tight targets, quick exits, small profits taken eagerly — because such systems feel good far more frequently than they feel bad. The trap, as Chapter One's data already showed, is that a high hit rate can comfortably coexist with a system that loses money, once the size of the rare loss is large enough to offset many small wins.

EXAMPLE — A 90% hit rate that goes nowhere

Consider a system that risks ₹900 to make ₹100 — a 9:1 risk-to-reward ratio in the wrong direction — but wins 90% of the time. $EV = 0.9 \times (+₹100) + 0.1 \times (-₹900) = ₹90 - ₹90 = ₹0$. The system is breakeven before a single rupee of cost is added, despite winning nine trades out of every ten. Add the same ₹100 flat cost used throughout this notebook, and it turns reliably negative.

The lesson is not that hit rate is meaningless — it clearly is one half of the expected-value equation. The lesson, backed now by both theory and the FXCM evidence, is that **hit rate alone, without its paired ratio, tells a trader almost nothing about whether a system is actually profitable.**

Chapter Four

The Ceiling — Why 50% Is Already Near the Top

A second, related instinct is to imagine that hit rate can be pushed indefinitely higher through better skill or sharper timing — 60%, 70%, even 80% consistently. For most short-duration directional bets, including intraday BankNifty option trades taken on a technical setup, this expectation does not hold up, and it is worth being explicit about why — and the FXCM data in Chapter One already hints at the answer: even traders sitting above 60% still lost, so a higher hit rate was never the missing ingredient.

A short-term directional trade is, at its core, a bet on which way price moves next over a fairly small window. Over such short windows, price action in a liquid index behaves close to a random walk — closer to a coin toss than most traders are comfortable admitting. A genuine, skill-based edge in reading price action can nudge the true win probability a few points above a pure 50/50 split, but it rarely moves it very far. In practice, most durable directional systems settle somewhere in a band of roughly 40% to 50% — genuinely better than a coin flip in expectation, but nowhere near a coin flip in appearance.

This is precisely why the profit ratio exists as a lever in the first place. If hit rate is structurally capped near the 50% neighborhood, a system that only wins when its hit rate is unusually high is a system quietly betting against how markets actually behave over short windows — the same trap Chapter One's traders fell into.

EXAMPLE — The breakeven hit rate this notebook is actually running at

For the core 2:1 setup used throughout this notebook (net win ₹900, net loss ₹600, after cost), the breakeven hit rate is found by solving $p \times 900 = (1-p) \times 600$, which gives $p = 600 / 1500 = 40\%$. In other words, this exact system stays profitable at a hit rate as low as 40% — squarely inside the realistic 40-to-50% band described above, and well below the assumed 50%, giving the plan a cushion that the FXCM traders never built in for themselves.

Seen this way, **the 50% hit rate used across this notebook is not a conservative placeholder — it is close to the realistic ceiling for this kind of trade**, and the system has been built with a comfortable cushion (40% breakeven versus a 50% assumed rate) precisely because hit rate itself cannot be relied upon to climb much further.

Chapter Five

What Actually Moves the Needle: Entries, Exits, Ratio

If hit rate is bounded by the nature of short-term price movement, then the lever actually within a trader's control is the profit ratio — and the profit ratio is not set by luck. It is set by the quality of entries and exits: how early a valid setup is recognized, how tightly a losing trade is cut, and how much of a winning move is allowed to run before booking profit. This is exactly the fix the FXCM research pointed to when it found ratio-disciplined traders profiting three times as often as the rest.

This is a fundamentally different skill from prediction. A trader does not need to know, in advance, which trades will win in order to improve their ratio — they need only to consistently exit losers early and let winners extend further. Improving entries and exits pushes the ratio; improving prediction pushes the hit rate. Given the ceiling described in Chapter Four, **the ratio is the more elastic of the two.**

EXAMPLE — The same hit rate, a different ratio, a very different month

This notebook already ran this exact comparison. At a 50% hit rate with a 1.5:1 ratio (win ₹750, lose ₹500, after cost), one trade a day over twenty days produced an expected ₹500 for the month. At a 2:1 ratio (win ₹1,000, same loss, same cost), the identical hit rate and identical trade count produced ₹3,000 — six times the result, from the same skill at calling direction.

The practical implication for a trader working within a realistic 40-to-50% hit rate band is this: **time spent trying to predict direction more often has a hard, structural ceiling. Time spent tightening stop discipline and extending winners has a much longer runway before it hits diminishing returns.**

The One Rule That Actually Matters

It is worth stating this as plainly as possible, because it is easy to nod along with the reasoning above and still quietly keep chasing hit rate in practice. **In the long run, your hit rate does not matter on its own.** Not 45%, not 55%, not 65%. A hit rate is not a measure of a system's health. It is only ever one half of a fraction, and a fraction with no denominator worth anything tells you nothing.

What actually matters, the only thing that matters over a long enough run, is **whether the system carries a positive expected value** once every trade's real payout, real

loss, and real cost are netted against each other. If that number is positive, a 40-to-50% hit rate is not a weakness to apologize for — it is entirely sufficient, exactly as Chapter Four's breakeven math showed. A trader sitting at 42% with a well-kept 2:1 ratio is in a mathematically stronger position than a trader sitting at 65% with a ratio that quietly runs against them, the way Chapter One's FXCM traders did.

This is why entries and exits, not prediction, deserve the discipline. **A trader's technical entries and exits should be built to secure a minimum profit ratio of 1:2** — risking one unit to target at least two — before hit rate is even part of the conversation. Get the ratio to that floor first, at whatever hit rate the setup honestly delivers, and the arithmetic in the rest of this notebook takes care of the outcome from there.

“Fix the ratio first. Let the hit rate be whatever the market honestly gives you — 40%, 50%, it does not matter, as long as the edge is positive.”

Chapter Six

Three Ways to Stop — The Conditional Ladder

The first daily structure proposed was a ladder with a built-in stopping rule: win the first trade and close for the day; lose it, and a second trade opens; lose that one too, and only then does a third trade open.

Working the tree outward from a single trade's 50/50 branch, four distinct daily paths emerge:

Path	Probability	Trades Taken	Day P&L
Win on Trade 1	50%	1	+₹900
Lose 1, Win 2	25%	2	+₹300
Lose 1 & 2, Win 3	12.5%	3	-₹300
Lose all three	12.5%	3	-₹1,800

Table 6.1 – Outcome tree for the stop-on-first-win ladder.

Daily Structure	Daily EV	20-Day EV	Win-Day Prob.	Worst Day
1 trade/day	+₹150	+₹3,000	50%	-₹600
2 trades/day (stop on win)	+₹225	+₹4,500	75%	-₹1,200
3 trades/day (stop after 2 wins)	+₹262.50	+₹5,250	87.5%	-₹1,800

Table 6.2 – The conditional ladder, one to three trades a day.

EXAMPLE — A concrete twenty days

Imagine this ladder run for exactly twenty trading days with three trades allowed. On roughly 10 of those days the trader wins the first trade and stops (+₹900 each, ₹9,000 total). On roughly 5 days the first trade loses but the second wins (+₹300 each, ₹1,500 total). On about 2–3 days both of the first two lose but the third wins (near breakeven, small net loss). On the remaining 2–3 days all three lose (-₹1,800 each). Add these together and the total lands close to the ₹5,250 expected figure in Table 6.2 — though any single real month will scatter around that center rather than hit it exactly.

A clear pattern already announces itself: **allowing more trades a day does not merely add variance, it adds expected profit and improves the odds of ending any given day in the green.**

Chapter Seven

What Happens When You Never Stop

The natural next question was to test what happens if the stopping rule is dropped altogether — if a trader with a daily quota simply takes every trade in that quota, win or lose, without closing early on a win.

Always Two Trades

Result	Probability	Day P&L
Win, Win	25%	+₹1,800
Win, Lose (either order)	50%	+₹300
Lose, Lose	25%	-₹1,200

Table 7.1 - Two independent trades, both always taken.

Daily EV = +₹300, or **+₹6,000** across twenty days.

Always Three Trades

Result	Probability	Day P&L
3 wins	12.5%	+₹2,700
2 wins, 1 loss	37.5%	+₹1,200
1 win, 2 losses	37.5%	-₹300
0 wins (3 losses)	12.5%	-₹1,800

Table 7.2 - Three independent trades, all always taken.

Daily EV = +₹450, or **+₹9,000** across twenty days.

EXAMPLE — Why stopping on a win costs money here

Compare two trades on the same day. Trade 1 wins (+₹900). Under a stop-on-win rule, the day ends there. Under an always-take rule, Trade 2 is still taken, and since it is independent of Trade 1, it still carries its own +₹150 edge on average. Skipping it does not reduce risk retroactively on Trade 1 — it simply forfeits a separate, unrelated positive-EV opportunity that happened to arrive on a good day.

Because every trade carries the identical, independent ₹150 edge, **expected profit scales with the raw count of trades taken — not with the order they arrive in, and not with whether earlier trades in the day won or lost.**

Chapter Eight

A Softer Edge — Testing a 1.5:1 Ratio

Before returning to the 2:1 ratio, the notebook paused to test a gentler profile described as easier to achieve in practice: a 1.5:1 profit ratio, still at a 50% hit rate, with the loss held at ₹500 and the profit target reduced to ₹750.

Outcome	Gross Result	Cost	Net Result
Win	₹750	-₹100	+₹650
Loss	-₹500	-₹100	-₹600

Table 8.1 – Net result of a single trade at 1.5:1.

EV(one trade) = +₹25 — roughly a sixth of the edge available at 2:1.

Daily Structure	Daily EV	20-Day EV	Worst Day
1 trade/day	+₹25	+₹500	-₹600
2/day (stop on win)	+₹37.50	+₹750	-₹1,200
3/day (stop after 2 wins)	+₹43.75	+₹875	-₹1,800
2/day (always both)	+₹50	+₹1,000	-₹1,200
3/day (always all three)	+₹75	+₹1,500	-₹1,800

Table 8.2 – The 1.5:1 ratio across every daily structure.

EXAMPLE — Same risk, thinner reward

Notice the worst single day at three trades is still -₹1,800 — identical to the 2:1 version in Chapter Six. Only the reward changed, from ₹5,250 down to ₹875 expected per month at that same structure. A trader accepting a softer ratio is not accepting less risk on a bad day — **they are accepting less pay for the same risk.**

“A weaker ratio does not make the storm smaller. It only makes the harvest thinner.”

Chapter Nine

Returning to 2:1, and Counting by the Month

With the softer ratio set aside, the notebook returned to the original 2:1 profile and assembled every structure examined so far into a single monthly ledger, with twenty trading days treated as one month.

Structure	Trades/Day	Daily EV	Trades/Mo.	Monthly EV	Worst Day	Win Days
1 trade/day	1	+₹150	20	+₹3,000	-₹600	50%
2/day (stop on win)	1-2	+₹225	≈30	+₹4,500	-₹1,200	75%
3/day (stop after 2W)	1-3	+₹262.50	≈35	+₹5,250	-₹1,800	87.5%
2/day (always both)	2	+₹300	40	+₹6,000	-₹1,200	75%
3/day (always all 3)	3	+₹450	60	+₹9,000	-₹1,800	87.5%

Table 9.1 – Full monthly ledger at the 2:1 ratio, two-lot cost basis.

For the conditional, stop-on-win structures, the number of trades taken in a month is not fixed, since some days end after one trade and others run the full quota. **The monthly EV describes an expected outcome around which real results will scatter — it is a center of gravity, not a promise.**

Chapter Ten

Lot Size and the Cost Floor

Every calculation so far has carried a flat ₹100 cost per trade for platform fees, brokerage, and slippage — the same charge whether the position is one lot or three. This flat structure has a quiet consequence worth making explicit: the smaller the position, the larger a bite that fixed cost takes out of the gross result, and the larger the position, the smaller that bite becomes as a share of the trade.

Assume the 2:1 ratio scales cleanly with lot size — each lot contributing ₹500 toward a win and ₹250 toward a loss — while the ₹100 cost stays fixed regardless of how many lots are traded:

Lot Size	Gross Win	Gross Loss	Cost	Net Win	Net Loss	Per-Trade EV	Cost % of Win
1 lot	₹500	-₹250	-₹100	+₹400	-₹350	+₹25	20.0%
2 lots	₹1,000	-₹500	-₹100	+₹900	-₹600	+₹150	10.0%
3 lots	₹1,500	-₹750	-₹100	+₹1,400	-₹850	+₹275	6.7%

Table 10.1 – Cost dilution as lot size grows, 2:1 ratio.

EXAMPLE — Doubling size more than doubles the edge

Going from 1 lot to 2 lots doubles the gross win and gross loss exactly, as expected. But because the ₹100 cost does not double along with it, the per-trade EV goes from ₹25 to ₹150 — six times larger, not two. The cost, which ate 20% of the gross win at 1 lot, only eats 10% at 2 lots. This is why the three-lot figures used from Chapter Seven onward (+₹275 per trade) are noticeably better than a simple 1.5× scaling of the two-lot figures (+₹150) would suggest — $1.5 \times 150 = ₹225$, not ₹275. The extra ₹50 comes purely from the fixed cost being spread over a larger position.

The practical implication is that **a trader operating near the minimum viable lot size is fighting cost drag on top of market risk** — the ₹100 charge is doing proportionally more damage the smaller the position gets. This argues for sizing up to whatever level a trader's capital and risk tolerance genuinely support, rather than trading the smallest size available.

This is not an argument for unlimited size. Larger positions bring their own frictions that this simple model does not capture — wider slippage on entry and exit in

BankNifty options as size grows, margin requirements that scale with lot count, and the psychological weight of a larger rupee swing on a losing day. The point of this chapter is narrower: **cost, specifically, becomes a smaller drag as size increases**, and that effect should be counted deliberately rather than ignored when comparing lot sizes.

Chapter Eleven

The Middle Path — A Conditional Third Trade

A more refined structure was proposed next: always take the first two trades of the day, regardless of outcome, and reserve the third trade for the single case where both of the first two were lost. Using the three-lot figures from Chapter Ten:

Path	Probability	Day P&L
Win, Win (no 3rd trade)	25%	+₹2,800
Win/Lose or Lose/Win (no 3rd)	50%	+₹550
Lose, Lose, then Win 3rd	12.5%	-₹300
Lose, Lose, then Lose 3rd	12.5%	-₹2,550

Table 11.1 – Two trades always, a third only after two losses.

Daily EV = **+₹618.75**, or **+₹12,375** across the month.

Structure	Daily EV	Monthly EV	Worst Day	Win-Days
1 trade/day	+₹275	+₹5,500	-₹850	50%
2/day (stop on win)	+₹412.50	+₹8,250	-₹1,700	75%
3/day (stop after 2 wins)	+₹481.25	+₹9,625	-₹2,550	87.5%
2/day always, 3rd only if both lose	+₹618.75	+₹12,375	-₹2,550	75%
3/day (always all three)	+₹825	+₹16,500	-₹2,550	87.5%

Table 11.2 – Where the conditional-third structure sits among the others.

EXAMPLE — Why the hybrid lands in the middle

On days that end Win-Win or a split, the hybrid behaves exactly like “always take two” — the third trade is never touched. Only on the roughly 1-in-8 days where both of the first two trades lose does the hybrid pick up the extra, positive-EV third trade. It therefore lands between “always two” (₹11,000) and “always three” (₹16,500).

Chapter Twelve

The Argument for No Ceiling

Under the exact assumptions given — a flat 50% hit rate, a fixed 2:1 ratio, and independence between trades — there is no natural stopping point at three. Every trade carries the same positive expected value regardless of what came before it, so the logic that favors a third trade over stopping at two applies with equal force to a fourth trade over a third, and a fifth over a fourth.

“The theoretical ceiling isn’t three — it is however many valid setups exist.”

This reframes the real constraint. **The question is not how many trades a trader should permit themselves in the abstract, but how many trades genuinely meet entry criteria on a given day.**

EXAMPLE — If five setups appeared instead of three

Following the same logic as Table 7.2, five independent trades always taken at the same three-lot, ₹275-per-trade edge would produce a daily EV of $5 \times ₹275 = ₹1,375$, or ₹27,500 across twenty days.

Two honest cautions belong here. First, real hit rates are rarely as flat and independent as assumed; a losing streak in a live market often signals a genuine shift in conditions rather than ordinary noise — exactly the kind of shift the FXCM data in Chapter One shows can quietly erode an otherwise decent hit rate. Second, a trade-count ceiling is a blunt instrument for managing risk; a rupee-based daily stop, decoupled from trade count, offers better protection without artificially capping upside on a good day.

Chapter Thirteen

Three Setups a Day — The Chosen Plan

With the ceiling question settled in principle, the practical constraint was supplied: in real trading conditions, only three genuinely valid setups tend to appear in a day. Given that constraint, **all three should be taken, every day, without a stop-on-win rule.**

Metric	Value
Daily EV	+₹825
Monthly EV (20 days)	+₹16,500
Worst single day	-₹2,550
Probability of a green day	87.5%
Probability of a full-loss day	12.5% ($\approx$ 2.5 days/month)

Table 13.1 - The adopted plan: three setups a day, all taken, three-lot sizing.

EXAMPLE — One real month, in shorthand

Across 20 trading days at 87.5% expected win-days: roughly 17-18 days close green (averaging well above ₹0 each), and roughly 2-3 days close as full-loss days at -₹1,800 to -₹2,550 each depending on lot size. Netting these together lands close to the ₹16,500 monthly figure — a plan that wins far more often than it loses, but whose rare bad days are meaningfully sized.

Chapter Fourteen

Guarding Against the Day the Math Breaks

Having settled on a structure with no intraday stopping rule, the notebook turns to a different kind of safeguard: one that watches across days for a signal that the founding assumption — a flat, independent 50% hit rate — may have quietly stopped being true, the same way it quietly stopped holding for the traders in Chapter One.

The Rule

Take all three trades every day as normal. If a day ends in a complete loss — all three trades lost — skip trading entirely on the following day as a deliberate cool-off, then resume the normal structure.

Event	Probability	Approx. Frequency
A single full-loss (LLL) day	12.5%	≈2.5 days / month
Two full-loss days in a row	1.56%	rare, not typical in most months

Table 14.1 - How often the trigger conditions actually occur.

Version	Trading Days/Mo.	Monthly EV	Worst Day
No pause (always three)	20	+₹16,500	-₹2,550
Pause one day after any LLL day	≈17.5	+₹14,440	-₹2,550
Recommended: pause after 2 consecutive LLL days	≈19.7	≈₹16,300	-₹2,550

Table 14.2 - Expected monthly cost of each safeguard.

EXAMPLE — Reading the two safeguards side by side

Pausing after every single full-loss day costs about ₹2,060 a month in forgone expected profit to guard against an event (12.5% likely) that is, under a genuinely flat 50% hit rate, ordinary noise. Pausing only after two full-loss days in a row guards against a much rarer event (1.56% likely) while giving up only about ₹200 of expected monthly profit — keeping roughly 99% of the theoretical ceiling intact.

“A single bad day is the price of the edge. Two bad days in a row is a question worth asking.”

Closing Note

What This Notebook Is, and Isn't

Everything in these pages follows logically from a handful of numbers held fixed: a 2:1 profit ratio, a 50% hit rate, a flat ₹100 cost per trade, and three genuine setups appearing in a day. Within that frame, the conclusion is unambiguous — **take every qualifying trade, size up enough to dilute the fixed cost, spend improvement effort on the ratio rather than the hit rate**, and reserve intervention for a genuine, rare break in the pattern rather than for ordinary bad luck.

But it is worth restating plainly what holds this entire structure up: the assumption that the 50% hit rate is real, flat, and independent of the day's sequence. **A real hit rate that drifts to 40% instead of 50% does not merely shrink these figures** — since 40% is this system's exact breakeven point, it can flip the entire plan from profitable to unprofitable outright.

The numbers in this book are therefore best read as a ceiling and a map — a clear picture of what a stated edge is worth if it holds, and a clear warning of how quickly that picture changes if it doesn't. Chapter One is a reminder that this warning is not academic: the same pattern this notebook warns against — a comfortable hit rate paired with a damaging ratio — is the documented, recurring failure mode of real retail traders at scale. **The discipline that matters most, in the end, is not the trade-count rule chosen on any given day, but the honesty to notice when the real world has quietly stopped matching the model — and the humility to keep the ratio, not the win rate, as the number worth defending.**

*Appendix***Full Reference Tables****A.1 - FXCM Study Summary**

Metric	Value
Trades closing at a profit (EUR/USD)	≈61%
Average winning trade	≈48 pips
Average losing trade	≈83 pips
Resulting average EV per trade	-3.09 pips
Profit rate with 1:1+ ratio discipline	≈53%
Profit rate without ratio discipline	≈17%

A.2 - Master Comparison, Two-Lot Sizing, 2:1 Ratio

Structure	Daily EV	Monthly EV	Worst Day	Win-Days
1 trade/day	+₹150	+₹3,000	-₹600	50%
2/day (stop on win)	+₹225	+₹4,500	-₹1,200	75%
3/day (stop after 2 wins)	+₹262.50	+₹5,250	-₹1,800	87.5%
2/day (always both)	+₹300	+₹6,000	-₹1,200	75%
3/day (always all three)	+₹450	+₹9,000	-₹1,800	87.5%

A.3 - Master Comparison, Three-Lot Sizing, 2:1 Ratio

Structure	Daily EV	Monthly EV	Worst Day	Win-Days
1 trade/day	+₹275	+₹5,500	-₹850	50%
2/day (stop on win)	+₹412.50	+₹8,250	-₹1,700	75%
3/day (stop after 2 wins)	+₹481.25	+₹9,625	-₹2,550	87.5%
2/day (always both)	+₹550	+₹11,000	-₹1,700	75%
2/day always, 3rd only if both lose	+₹618.75	+₹12,375	-₹2,550	75%
3/day (always all three)	+₹825	+₹16,500	-₹2,550	87.5%

A.4 - Comparison at a Softer 1.5:1 Ratio, Two-Lot Sizing

Structure	Daily EV	Monthly EV	Worst Day
1 trade/day	+₹25	+₹500	-₹600
2/day (stop on win)	+₹37.50	+₹750	-₹1,200
3/day (stop after 2 wins)	+₹43.75	+₹875	-₹1,800
2/day (always both)	+₹50	+₹1,000	-₹1,200
3/day (always all three)	+₹75	+₹1,500	-₹1,800

A.5 - Cost Dilution by Lot Size, 2:1 Ratio

Lot Size	Net Win	Net Loss	Per-Trade EV	Cost % of Gross Win
1 lot	+₹400	-₹350	+₹25	20.0%
2 lots	+₹900	-₹600	+₹150	10.0%
3 lots	+₹1,400	-₹850	+₹275	6.7%

A.6 - The Streak-Based Safeguard

Version	Trading Days/Mo.	Monthly EV	Worst Day
No pause (always three)	20	+₹16,500	-₹2,550
Pause one day after any LLL day	≈17.5	+₹14,440	-₹2,550
Recommended: pause after 2 consecutive LLL days	≈19.7	≈₹16,300	-₹2,550

All figures in this notebook are theoretical expectations derived from stated assumptions and do not constitute financial advice. Trading BankNifty options carries real risk of loss; past or modeled expectancy does not guarantee future results. The FXCM figures cited in Chapter One and Appendix A.1 are drawn from publicly published broker research and are summarized here in original wording, not reproduced verbatim.